

# Large Language Models

## Scaling Laws

ELL881 · AIL821



**Sourish Dasgupta**

Assistant Professor, DA-IICT, Gandhinagar

<https://daiict.ac.in/faculty/sourish-dasgupta>

# Kaplan Scaling Laws at a glance:

Parameters	Data	Compute	Batch Size	Equation
$N$	$\infty$	$\infty$	Fixed	$L(N) = (N_c/N)^{\alpha_N}$
$\infty$	$D$	Early Stop	Fixed	$L(D) = (D_c/D)^{\alpha_D}$
Optimal	$\infty$	$C$	Fixed	$L(C) = (C_c/C)^{\alpha_C}$ (naive)
$N_{\text{opt}}$	$D_{\text{opt}}$	$C_{\text{min}}$	$B \ll B_{\text{crit}}$	$L(C_{\text{min}}) = (C_c^{\text{min}}/C_{\text{min}})^{\alpha_C^{\text{min}}}$
$N$	$D$	Early Stop	Fixed	$L(N, D) = \left[ \left( \frac{N_c}{N} \right)^{\frac{\alpha_N}{\alpha_D}} + \frac{D_c}{D} \right]^{\alpha_D}$
$N$	$\infty$	$S$ steps	$B$	$L(N, S) = \left( \frac{N_c}{N} \right)^{\alpha_N} + \left( \frac{S_c}{S_{\text{min}}(S, B)} \right)^{\alpha_S}$

Power Law	Scale (tokenization-dependent)
$\alpha_N = 0.076$	$N_c = 8.8 \times 10^{13}$ params (non-embed)
$\alpha_D = 0.095$	$D_c = 5.4 \times 10^{13}$ tokens
$\alpha_C = 0.057$	$C_c = 1.6 \times 10^7$ PF-days
$\alpha_C^{\text{min}} = 0.050$	$C_c^{\text{min}} = 3.1 \times 10^8$ PF-days
$\alpha_B = 0.21$	$B_* = 2.1 \times 10^8$ tokens
$\alpha_S = 0.76$	$S_c = 2.1 \times 10^3$ steps



# Is there any other alternative law?



# Turns out there is!

$$\begin{aligned}\text{Loss}(N_T, D) &= \frac{N_c}{N_T^\alpha} + \frac{D_c}{D^\beta} + \underbrace{E}_{\text{Lower bound}}, \\ \text{Loss}(N_T, C_T) &= \frac{N_c}{N_T^\alpha} + \frac{D_c}{(C_T/6N_T)^\beta} + E\end{aligned}$$

**Chinchilla (Hoffman) Scaling Law**



# The Chinchilla (Hoffman) Scaling Law

$$\text{Loss}(N_T, D) = \frac{N_c}{N_T^\alpha} + \frac{D_c}{D^\beta} + E$$

$$L(N, D) = 1.69 + \frac{406.4}{N^{0.34}} + \frac{410.7}{D^{0.28}}$$

$$N_{\text{opt}}(C) = G(C/6)^a \quad D_{\text{opt}}(C) = G^{-1}(C/6)^b$$

where  $G = \left(\frac{\alpha A}{\beta B}\right)^{\frac{1}{\alpha+\beta}}$      $a = \frac{\beta}{\alpha + \beta}$      $b = \frac{\alpha}{\alpha + \beta}$

Fitting the constants, yields:  $\alpha \approx \beta$   
i.e. equal scaling of **N** and **D**.



# Chinchilla Scaling Law vs. Kaplan Scaling Law

$$\text{Kaplan: } N_{\setminus E}^* \propto C_{\setminus E}^{0.73}$$

$$\text{Chinchilla: } N_T^* \propto C_T^{0.50}.$$

Performance penalty is  $N^{0.75} / D$

- if model increases 8x, dataset must increase 5x

VS.

Fitting the constants, yields:  $\alpha \approx \beta$   
i.e. equal scaling of  $N$  and  $D$ .

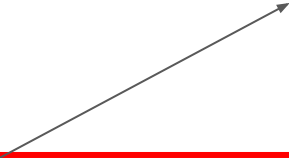
$$N_T = N_E + N_{\setminus E}, \quad C_T = 6N_T D = 6(N_E + N_{\setminus E})D,$$

$$N_E = (h + v)d, \quad C_{\setminus E} = 6N_{\setminus E}D.$$




# The (revised) Chinchilla Scaling Law

$$L(N, D) = 1.82 + \frac{514.0}{N^{0.35}} + \frac{2115.2}{D^{0.37}}$$


$$L(N, D) = 1.69 + \frac{406.4}{N^{0.34}} + \frac{410.7}{D^{0.28}}$$



# Is it a problem with our point-of-view?






# LLMs “seems” to get more intelligent with the following:



Amount of training data



Model size (parameters)



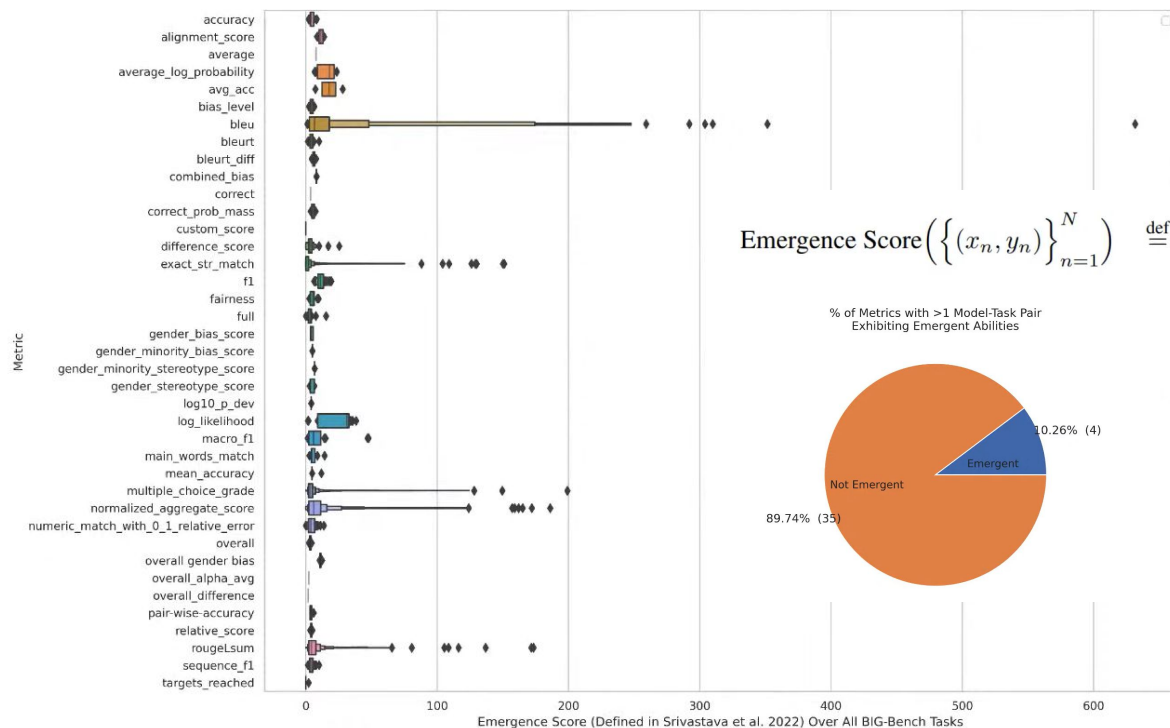
Amount of compute (or time)

$$\mathcal{L}_{CE}(N) = \left(\frac{N}{c}\right)^{\alpha}$$






# Motivation: Not all metrics score same (Emergence Score)



$$\text{Emergence Score} \left( \left\{ (x_n, y_n) \right\}_{n=1}^N \right) \stackrel{\text{def}}{=} \frac{\text{sign}(\arg \max_i y_i - \arg \min_i y_i) (\max_i y_i - \min_i y_i)}{\sqrt{\text{Median}(\{(y_i - y_{i-1})^2\}_i)}}$$






# Is your accuracy metric non-linear or discontinuous?



> 92% of BIG-BENCH:

Multiple Choice Grade  $\stackrel{\text{def}}{=} \begin{cases} 1 & \text{if highest probability mass on correct option} \\ 0 & \text{otherwise} \end{cases}$

Exact String Match  $\stackrel{\text{def}}{=} \begin{cases} 1 & \text{if output string exactly matches target string} \\ 0 & \text{otherwise} \end{cases}$




*Too challenging for smaller models!  
Is it really worth??*





# Power Law in play!




$$\mathcal{L}_{CE}(N) = \left(\frac{N}{c}\right)^{\alpha} = -\log \hat{p}_{v^*}(N)$$

$$\hat{p}_{v^*}(N) = \exp\left(-\left(N/c\right)^{\alpha}\right)$$

$$L(D_0, N) = B_0 + \frac{B_1}{N^{\alpha_w}} + \dots$$


Length of token





# Problem with Non-linear Measure: Eg.: Exact string match

**Task:** Add  $k$ -digit integers



- 1 if all  $K+1$  digits in model's output are correct
- 0 otherwise

$$\hat{p}_{v^*}(N) = \exp\left(- (N/c)^\alpha\right)$$


$$\text{Accuracy}(N) \approx p_N(\text{single token correct})^{\text{num. of tokens}}$$





# Change of perspective: Measure: Edit distance

**Task:** Add  $k$ -digit integers



- 1 if all  $K+1$  digits in model's output are correct
- 0 otherwise

$$\hat{p}_{v^*}(N) = \exp\left(- (N/c)^\alpha\right)$$




$$\text{Edit Distance}(N) \approx L \left(1 - p_N(\text{single token correct})\right)$$





# Problem with Discontinuous Measure: Eg.: MCG

**Task:** Choose one of two



- 1 if highest probability mass on correct option
- 0 otherwise


$$\hat{p}_{v^*}(N) = \exp\left(- (N/c)^\alpha\right)$$





# Change of perspective: Measure: Brier Score

**Task:** Choose one of two



- 1 if all  $K+1$  digits in model's output are correct
- 0 otherwise

$$\hat{p}_{v^*}(N) = \exp\left(- (N/c)^\alpha\right)$$





Brier Score =  $(1 - \text{probability mass on correct option})^2$







# Prediction: Power Law vs. Near-Linear counterpart




$$\hat{p}_{v^*}(N) = \exp\left(- (N/c)^\alpha\right)$$






# Results on GPT3.5/3: Task: 2-digit integer multiplication




$$\hat{p}_v^*(N) = \exp\left(- (N/c)^\alpha\right)$$





# Does the claim work for Google BIG-BENCH benchmark?





# Key Takeaways

- Want to predict **without the theatrics**? Choose a *metric that's "soft"* (in the continuous sense)
- There's *no sudden jump* in reality ("*most*" can be predicted on a near-linear scale)
- Do we really need the power law of scale? Maybe not!

